Worth Using

Live pilot and full-run budget

BLOCKED BEFORE LIVE CALLS
No efficacy conclusion is drawn from this pilot. Provider usage, accepted settings and cost remain unmeasured.

Source commit: 8a3a3ac5c26b007d592c0f423f74b52c157e00e9

Model requests and provider responses

A successful request confirms payload acceptance. Temperature and token limits are request values, not server-echoed effective settings.

Role	Requested model	Exact returned model	Success / error
eval	UNSET	None returned	0 / 0
judge	UNSET	None returned	0 / 0
simulator	UNSET	None returned	0 / 0

Token cost per conversation

No live conversation exists. Token usage, latency and dollar cost are unavailable.

The planned full run

Development	Private held-out	Separate pilot
12 fixtures · 3 seeds 324 conversations	4 fixtures · 3 seeds 108 conversations	2 fixtures · 1 seed 18 conversations

Combined main plan: 432 conversations, 864 blind judgments, 48 simulator generations and 48 script audits. Private fixture loading is not yet integrated into the current runner. All private payloads remain outside the public repository.

No pilot-based full-run cost projection is available. No model price or token quantity has been invented.

Errors and next gate

preflight: EVAL_MODEL is unset

preflight: OPENAI_API_KEY is unset

preflight: JUDGE_MODEL is unset

preflight: SIMULATOR_MODEL is unset (defaults to EVAL_MODEL)

preflight: EVAL_MODEL_FAMILY is unset; declare the actual model family explicitly

preflight: JUDGE_MODEL_FAMILY is unset; declare the actual model family explicitly

Stop before the full run. The user must receive the pilot measurements and projected cost first. Full-run headline comparisons must hold on the private cohort and show rating coverage, evidence coverage and evidence recognition next to unsupported upgrades. Missing ratings cannot masquerade as honesty.

No efficacy intervals or illustrative wins/losses are claimed here. Final excerpts must use real full-run transcripts. Fictional fixtures, fixed simulated users, one repeated judge and a small held-out sample limit generalization. Real tools with real usage logs belong to a later separate cohort.

Audit files: REPORT.md; scores/pilot_conversations.csv; scores/pilot_cost_latency.csv; scores/pilot_accepted_settings.json; scores/pilot_projection.json. Live requests appear in pilot/live/calls.jsonl only when calls exist.